

SQL, the Database, and You

John Mosesman

Developer, consultant, database-liker



Why this talk?

- Your app is your data
- Databases are one of the first and almost universal dependencies
- Databases can be one of the first performance bottlenecks
- Data storage and querying is not very intuitive
- Useful knowledge for all types of devs

Talk structure

1. Structuring data for a relational database
2. Querying (SQL)
3. Associations
4. Data integrity
5. Performance

1. Structuring data for a relational database



petco

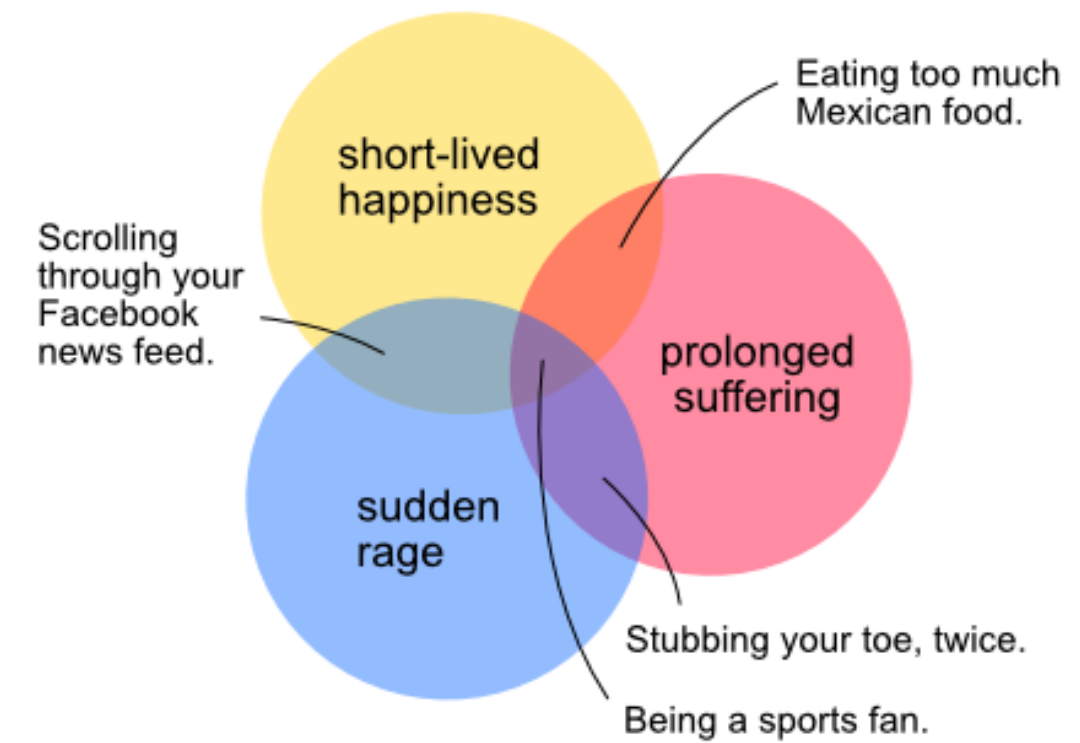
grooming

FIRE LANE
NO
PARKING
TOW-AWAY ZONE
CITY OF LAS VEGAS

101




```
{
  person: {
    name: "John Smith",
    age: 50,
    email: "john@smith.com"
  }
}
```



The following data were generated randomly for demonstration purposes only. Any resemblance to real people or places is purely coincidental.							
1	First Name	Last Name	Raffle #	City	State	ZIP	Street
2	Marlene	Culbertson	15121	Collierville	TN	38017	3550
3	John	Westbrook	15443	New York	NY	10022	2700
4	Eli	Downs	12677	Quincy	MA	2170	92 L
5	Donald	Pierce	15254	Kingsland	TX	78639	1420
6	Salvador	Lantz	14811	Gilbert	AZ	85297	6000
7	David	Mccall	12488	New york	TX	77024	8 St
8	Tim	Browder	14456	Tuscaloosa	IL	60041	2700
9	Barbara	Culbertson	13905	Waxhaw	NC	2/17/1977	70 L
10	Steve	Tate	14059	Gorham	ME	4038	9400
11	Diane	Freeman	12647	Tuscaloosa	CA	92024	2000
12	Frans	Teklinski	12907	Kirkwood	MO	63122	1510
13	Joshua	Langston	13547	Bellingham	WA	98225	27 S
14	Alan	Rose	15187	Richmond	VA	23227	9420
15	It	Schrimpf	12532	Clermont	FL	34711	5150
16	Christopher	Scott	15513	Manawa	WI	54949	4000
17	Stephen	Stearns	14657	Washington	DC	20016	4200
18	David	Oberg	13467	Denver	CO	80209	9000
19	Laurence	Longoria	15473	Powhatan	VA	23139	55 S

Random generator

Integer

From: 12343 To: 15550

☒ Unique values

Real

Boolean

Date

Strings

GENERATE

AbleBits

Databases

+-----+		+-----+		+-----+		+-----+	
food		tastiness		cost		health	
+-----+		+-----+		+-----+		+-----+	
sushi		10		10		7	
pizza		8		5		3	
salad		2		8		9	
water		0		0		10	
+-----+		+-----+		+-----+		+-----+	

What is "good" structure?

- Goal: to reduce data redundancy and improve data integrity
- Free from change anomalies

FEATURE REQUEST

"People and their pets"

"People and their pets"

person_name	person_age	work_place	work_phone	pet_name	pet_type	pet_legs
John	26	Pizza Palace	555-555-5555	Ruby	Dog	4
Susan	22	Pizza Hut	111-111-1111	Pete	Bird	2
Jim	30	Pizza Town	333-333-3333	Fluffy	Cat	4

Adding a pet

person_name	person_age	work_place	work_phone	pet_name	pet_type	pet_legs
John	26	Pizza Palace	555-555-5555	Ruby	Dog	4
Susan	22	Pizza Hut	111-111-1111	Pete	Bird	2
Jim	30	Pizza Town	333-333-3333	Fluffy	Cat	4
John	26	Pizza Palace	555-555-5555	Buddy	Dog	4

Adding a pet

person_name	person_age	work_place	work_phone	pet1_name	pet1_type	pet1_legs	pet2_name	pet2_type	pet2_legs
John	26	Pizza Palace	555-555-5555	Ruby	Dog	4	Buddy	Dog	4
Susan	22	Pizza Hut	111-111-1111	Pete	Bird	2			
Jim	30	Pizza Town	333-333-3333	Fluffy	Cat	4			

Owner gives up a pet

person_name	person_age	work_place	work_phone	pet_name	pet_type	pet_legs
John	26	Pizza Palace	555-555-5555	Ruby	Dog	4
???	??	Pizza Hut	111-111-1111	Pete	Bird	2
Jim	30	Pizza Town	333-333-3333	Fluffy	Cat	4

Database "Normalization"

The Three "Normal Forms"

- Introduced by Dr. Edgar F. Codd
- Meets our goal of "reduce data redundancy and improve data integrity"

1st Normal Form

"The information is stored in a relational table and each column contains atomic values, and there are not repeating groups of columns."¹

person_name	person_age	pet1_name	pet1_type	pet1_legs	pet2_name	pet2_type	pet2_legs
John	26	Ruby	Dog	4	Buddy	Dog	4
Susan	22	Pete	Bird	2			
Jim	30	Fluffy	Cat	4			

¹ <https://www.essentialsql.com/get-ready-to-learn-sql-database-normalization-explained-in-simple-english/>

1st Normal Form

people

person_name	person_age	work_place	work_phone
John	26	Pizza Palace	555-555-5555
Susan	22	Pizza Hut	111-111-1111
Jim	30	Pizza Town	333-333-3333

pets

person_name	pet_name	pet_type	pet_legs
John	Ruby	Dog	4
Susan	Pete	Bird	2
Jim	Fluffy	Cat	4

2nd Normal Form

"The table is in first normal form and all the columns depend on the table's primary key."¹

people

person_name	person_age	work_place	work_phone
John	26	Pizza Palace	555-555-5555
Susan	22	Pizza Hut	111-111-1111
Jim	30	Pizza Town	333-333-3333

¹ <https://www.essentialsql.com/get-ready-to-learn-sql-database-normalization-explained-in-simple-english/>

2nd Normal Form

people

person_name	person_age
John	26
Susan	22
Jim	30

work_places

person_name	work_place	work_phone
John	Pizza Palace	555-555-5555
Susan	Pizza Hut	111-111-1111
Jim	Pizza Town	333-333-3333

3rd Normal Form

"The table is in second normal form and all of its columns are not transitively dependent on the primary key."¹

Or

"Every non-key attribute must provide a fact about the key, the whole key, and nothing but the key, so help me Codd."²

¹ <https://www.essentialsql.com/get-ready-to-learn-sql-database-normalization-explained-in-simple-english/>

² Credited to Bill Kent on Wikipedia

3rd Normal Form

pets

person_name	pet_name	pet_type	pet_legs
John	Ruby	Dog	4
Susan	Pete	Bird	2
Jim	Fluffy	Cat	4
John	Buddy	Dog	4

3rd Normal Form

pets

person_name	pet_name	pet_type
John	Ruby	Dog
Susan	Pete	Bird
Jim	Fluffy	Cat
John	Buddy	Dog

pet_types

pet_type	pet_legs
Dog	4
Bird	2
Cat	4

Almost There

people

+-----+-----+	
person_name	person_age
+-----+-----+	
John	26
Susan	22
Jim	30
+-----+-----+	

pets

+-----+-----+-----+		
person_name	pet_name	pet_type
+-----+-----+-----+		
John	Ruby	Dog
Susan	Pete	Bird
Jim	Fluffy	Cat
John	Buddy	Dog
+-----+-----+-----+		

IDs and Primary Keys

Primary key - one or more columns that uniquely identifies each row

people

id	person_name	person_age
1	John	26
2	Susan	22
3	Jim	30

pets

id	person_id	pet_name	pet_type
20	1	Ruby	Dog
21	2	Pete	Bird
22	3	Fluffy	Cat
23	1	Buddy	Dog

Foreign Keys

Foreign key - a reference to a primary key in another table

people

id (PK)	person_name	person_age
1	John	26
2	Susan	22
3	Jim	30

pets

id (PK)	person_id (FK)	pet_name	pet_type
20	1	Ruby	Dog
21	2	Pete	Bird
22	3	Fluffy	Cat
23	1	Buddy	Dog

2. Querying

Structured Query Language (SQL)

Most modern Database Management Systems (DBMS) adhere to the SQL Standard.

Creating a table

```
CREATE TABLE people (  
    first_name VARCHAR,  
    last_name  VARCHAR,  
    age        INTEGER  
)
```

```
CREATE TABLE
```

Data types

- CHAR/VARCHAR/TEXT
- INTEGER
- BOOLEAN
- DATE/DATETIME

SELECT

```
SELECT *  
FROM people
```

0 rows

(* = all columns)

SELECT

```
SELECT first_name, last_name  
FROM people
```

0 rows

INSERT

```
INSERT INTO people (first_name, last_name, age) VALUES  
    ('John', 'Mosesman', 26),  
    ('Bob', 'Smith', 40),  
    ('Jane', 'Doe', 32)
```

```
INSERT 0 3
```

WHERE

```
SELECT *  
FROM people  
WHERE first_name = 'John'
```

+	-----	+	-----	+	-----	+
	first_name		last_name		age	
+	-----	+	-----	+	-----	+
	John		Mosesman		26	
+	-----	+	-----	+	-----	+

WHERE

```
INSERT INTO people  
SELECT 'John', 'Smith', 100
```

```
SELECT *  
FROM people  
WHERE first_name = 'John'
```

first_name	last_name	age
John	Mosesman	26
John	Smith	100

AND and OR

```
SELECT *  
FROM people  
WHERE first_name = 'John'  
      AND last_name = 'Mosesman'
```

```
SELECT *  
FROM people  
WHERE age > 50  
      OR age < 20
```


ALTER

```
ALTER TABLE people ADD COLUMN id SERIAL PRIMARY KEY;
```

```
---
```

```
SELECT *  
FROM people  
WHERE id = 1
```

```
---
```

id	first_name	last_name	age
1	John	Mosesman	26

UPDATE

```
UPDATE people  
SET first_name = 'Jack'  
WHERE id = 2
```

```
UPDATE 1
```

Aggregations

```
SELECT AVG(age) FROM people      -- 49.5
```

```
SELECT SUM(age) FROM people      -- 198
```

```
SELECT COUNT(*) FROM people      -- 4
```

3. *Associations*

Our Pet Store

people

id	name	age
1	John	26
2	Susan	22
3	Jim	30

pets

id	person_id	name	type	age
20	1	Ruby	Dog	3
21	2	Pete	Bird	5
22	3	Fluffy	Cat	12
23	1	Buddy	Dog	1

Our Pet Store (in SQL)

```
CREATE TABLE people (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR,  
  age INTEGER  
)
```

Our Pet Store (in SQL)

```
CREATE TABLE people (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR,  
    age INTEGER  
)
```

```
CREATE TABLE pets (  
    id SERIAL PRIMARY KEY,  
    person_id INTEGER REFERENCES people,  
    name VARCHAR,  
    age INTEGER  
)
```

FEATURE REQUEST

"Find all of a person's pets"

Finding pets

```
SELECT *  
FROM pets  
WHERE person_id = 1
```

id	person_id	name	age
1	1	Ruby	1
4	1	Buddy	3

FEATURE REQUEST

"Get everyone and their pets"

JOINS

"Join" two tables together by matching related rows.

people.id	people.name	people.age	pets.id	pets.person_id	pets.name	pets.age
1	John	26	20	1	Ruby	3
...						

Naive JOIN implementation

1. Combine every row from people with every row from pets
2. Filter out rows that don't match

Cartesian Product or "Cross Join"

letter	
A	
B	

number	
1	
2	

Cartesian Product or "Cross Join"

+-----+-----+	
letter	number
+-----+-----+	
A	1
A	2
B	1
B	2
+-----+-----+	

CROSS JOIN of people and pets

```
SELECT *
FROM people
  CROSS JOIN pets
```

16 rows

id	name	age	id	person_id	name	age
1	John Mosesman	26	1	1	Ruby	1
1	John Mosesman	26	2	2	Pete	5
1	John Mosesman	26	3	3	Fluffy	12
1	John Mosesman	26	4	1	Buddy	3
2	Bob Smith	40	1	1	Ruby	1
2	Bob Smith	40	2	2	Pete	5
2	Bob Smith	40	3	3	Fluffy	12
2	Bob Smith	40	4	1	Buddy	3
3	Jane Doe	32	1	1	Ruby	1
...						

CROSS JOIN of people and pets

id	name	age	id	person_id	name	age
1	John Mosesman	26	1	1	Ruby	1
1	John Mosesman	26	2	2 (???)	Pete	5
...						

CROSS JOIN with filter

```
SELECT *  
FROM people  
    CROSS JOIN pets  
WHERE people.id = pets.person_id
```

4 rows

id	name	age	id	person_id	name	age
1	John Mosesman	26	1	1	Ruby	1
2	Bob Smith	40	2	2	Pete	5
3	Jane Doe	32	3	3	Fluffy	12
1	John Mosesman	26	4	1	Buddy	3

JOIN

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id
```

JOIN filtering

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
        AND pets.age > 2
```

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
        AND pets.age > 2  
WHERE people.age < 30
```

Order of Operations

1. FROM

2. JOIN

3. WHERE

4. SELECT

Order of Operations

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
        AND pets.age > 2  
WHERE people.age < 30
```

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
WHERE people.age < 30  
    AND pets.age > 2
```

FEATURE REQUEST

"Track pet shot records."

Multiple JOINS

```
CREATE TABLE shot_records (  
    id SERIAL PRIMARY KEY,  
    pet_id INTEGER REFERENCES pets,  
    shot_name VARCHAR,  
    date DATE  
);
```

```
INSERT INTO shot_records (pet_id, shot_name, date) VALUES  
(1, 'Rabies', CURRENT_DATE),  
(1, 'Bordetella', CURRENT_DATE)
```

Multiple JOINS


```
SELECT *
FROM people
  JOIN pets ON pets.person_id = people.id
  JOIN shot_records sr ON sr.pet_id = pets.id
```

id	name	age	id	person_id	name	age	id	pet_id	shot_name	date
1	John Mosesman	26	1	1	Ruby	1	1	1	Rabies	2017-07-20
1	John Mosesman	26	1	1	Ruby	1	2	1	Bordetella	2017-07-20

Join Types


CROSS JOIN (Union)⁴

Full Join



⁴ <http://www.sql-join.com/sql-join-types/>


INNER JOIN (Intersection)⁴




⁴ <http://www.sql-join.com/sql-join-types/>

OUTER JOIN⁴

Left Join




Right Join




⁴ <http://www.sql-join.com/sql-join-types/>

Join Types⁴


Full Join




Left Join



Inner Join



Right Join



⁴ <http://www.sql-join.com/sql-join-types/>

A brief detour: NULL

NULL In Other Languages

```
irb(main):001:0> nil == nil  
=> true
```

```
irb(main):002:0> nil == 1  
=> false
```

NULL - "An unknown value"

```
SELECT 1 = 1
```


NULL - "An unknown value"

```
SELECT 1 = 1 -- TRUE
```

NULL - "An unknown value"

```
SELECT 1 = 1                -- TRUE
SELECT 1 <> 1
```

NULL - "An unknown value"

```
SELECT 1 = 1          -- TRUE
SELECT 1 <> 1         -- FALSE
```

NULL - "An unknown value"

```
SELECT 1 = 1                -- TRUE
SELECT 1 <> 1                -- FALSE
SELECT 1 = NULL
```

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	-- NULL

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	-- NULL
SELECT ' ' = NULL	

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	-- NULL
SELECT ' ' = NULL	-- NULL

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	-- NULL
SELECT ' ' = NULL	-- NULL
SELECT NULL = NULL	

NULL - "An unknown value"

SELECT 1 = 1	-- TRUE
SELECT 1 <> 1	-- FALSE
SELECT 1 = NULL	-- NULL
SELECT 0 = NULL	-- NULL
SELECT ' ' = NULL	-- NULL
SELECT NULL = NULL	-- NULL

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL -- TRUE
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL -- TRUE
SELECT NULL IS NULL
```


IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL -- TRUE
SELECT NULL IS NULL       -- TRUE
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL -- TRUE
SELECT NULL IS NULL       -- TRUE
SELECT NULL IS NOT NULL
```

IS NULL / IS NOT NULL

```
SELECT 1 IS NULL          -- FALSE
SELECT 'pizza' IS NOT NULL -- TRUE
SELECT NULL IS NULL       -- TRUE
SELECT NULL IS NOT NULL   -- FALSE
```

Back to JOINS

FEATURE REQUEST

"All people in our database, and their pet data
—if they have any."

OUTER JOIN (LEFT or RIGHT)

people LEFT JOIN pets => All people data, and pets if they have them

people RIGHT JOIN pets => All pets data, and people data if they have an owner

people LEFT JOIN pets

```
SELECT *
FROM people
  LEFT JOIN pets ON pets.person_id = people.id
```

id	name	age	id	person_id	name	age
1	John Mosesman	26	1	1	Ruby	1
2	Bob Smith	40	2	2	Pete	5
3	Jane Doe	32	3	3	Fluffy	12
1	John Mosesman	26	4	1	Buddy	3
4	Jack Smith	100				

people LEFT JOIN pets

```
SELECT *
FROM people
  LEFT JOIN pets ON pets.person_id = people.id
```

id	name	age	id	person_id	name	age
1	John Mosesman	26	1	1	Ruby	1
2	Bob Smith	40	2	2	Pete	5
3	Jane Doe	32	3	3	Fluffy	12
1	John Mosesman	26	4	1	Buddy	3
4	Jack Smith	100	NULL	NULL	NULL	NULL

People without pets

```
SELECT *  
FROM people  
    LEFT JOIN pets ON pets.person_id = people.id  
WHERE pets.id IS NULL
```

id	name	age	id	person_id	name	age
4	Jack Smith	100				

FEATURE REQUEST

"People with pets that don't have their shots."

Pets Without Shots

```
SELECT *
FROM people
  JOIN pets ON pets.person_id = people.id
  LEFT JOIN shot_records sr ON sr.pet_id = pets.id
WHERE sr.id IS NULL
```

id	name	age	id	person_id	name	age	id	pet_id	shot_name	date
2	Bob Smith	40	2	2	Pete	5				
1	John Mosesman	26	4	1	Buddy	3				
3	Jane Doe	32	3	3	Fluffy	12				

Pets Without Shots

```
SELECT people.*, pets.*
FROM people
  JOIN pets ON pets.person_id = people.id
  LEFT JOIN shot_records sr ON sr.pet_id = pets.id
WHERE sr.id IS NULL
```

id	name	age	id	person_id	name	age
2	Bob Smith	40	2	2	Pete	5
1	John Mosesman	26	4	1	Buddy	3
3	Jane Doe	32	3	3	Fluffy	12

FEATURE REQUEST

"Find each person and the number of pets they have."

Basic aggregations

```
SELECT AVG(age) FROM people
```

```
SELECT SUM(age) FROM people
```

```
SELECT COUNT(*) FROM people
```

Group by person

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id
```

Group by person

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
GROUP BY people.id
```


Group by person

```
SELECT *  
FROM people  
    JOIN pets ON pets.person_id = people.id  
GROUP BY people.id
```

```
ERROR:  column "pets.id" must appear in the GROUP BY clause or be used in an aggregate function  
LINE 1: SELECT *  
           ^
```

Group by person and get pet count

```
SELECT people.id, COUNT(pets.id)
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id
```

+	-	-	-	-	+	-	-	-	-	-	-	+
	id		count									
+	-	-	-	-	+	-	-	-	-	-	-	+
	2		1									
	1		2									
	3		1									
+	-	-	-	-	+	-	-	-	-	-	-	+

Group by person and get pet count

```
SELECT people.id, people.name, COUNT(pets.id)
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
```

id	name	count
2	Bob Smith	1
1	John Mosesman	2
3	Jane Doe	1

FEATURE REQUEST

"People with more pets are more likely to adopt more pets."

ORDER BY

```
SELECT people.id, people.name, COUNT(pets.id)
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
ORDER BY COUNT(pets.id) DESC
```

id	name	count
1	John Mosesman	2
2	Bob Smith	1
3	Jane Doe	1

ALIAS

```
SELECT people.id, people.name, COUNT(pets.id) AS num_pets
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
ORDER BY num_pets DESC
```

FEATURE REQUEST

"People with 2 or more pets are more likely to adopt more pets."

HAVING

```
SELECT people.id, people.name, COUNT(pets.id) AS num_pets
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
HAVING num_pets >= 2
ORDER BY num_pets DESC
```


HAVING

```
SELECT people.id, people.name, COUNT(pets.id) AS num_pets
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
HAVING num_pets >= 2
ORDER BY num_pets DESC
```

```
ERROR:  column "num_pets" does not exist
LINE 5: HAVING num_pets >= 2
              ^
```

Order of Operations, Again.

1. FROM
2. ON (JOIN)
3. WHERE
4. GROUP BY
5. HAVING
6. SELECT <-- ALIAS happens here
7. ORDER BY <-- Can be used here as well

HAVING, Again.

```
SELECT people.id, people.name, COUNT(pets.id) AS num_pets
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
HAVING COUNT(pets.id) >= 2
ORDER BY num_pets DESC
```

id	name	num_pets
1	John Mosesman	2

WHERE vs HAVING

"HAVING eliminates group rows that do not satisfy the condition. HAVING is different from WHERE: WHERE filters individual rows before the application of GROUP BY, while HAVING filters group rows created by GROUP BY."³

³ Postgres docs - <https://www.postgresql.org/docs/9.0/static/tutorial-agg.html>

LIMIT

```
SELECT people.id, people.name, COUNT(pets.id) AS num_pets
FROM people
      JOIN pets ON pets.person_id = people.id
GROUP BY people.id, people.name
HAVING COUNT(pets.id) >= 2
ORDER BY num_pets DESC
LIMIT 5
```

"Has Many"

- A person "has many" pets (one-to-many relationship)
- A pet "belongs to" an owner (one-to-one relationship).

FEATURE REQUEST

"Pets can belong to more than one owner."

Has And Belongs To Many (HABTM)

```
CREATE TABLE people (  
  id SERIAL PRIMARY KEY,  
  ...  
);
```

```
CREATE TABLE pets (  
  id SERIAL PRIMARY KEY,  
  ...  
)
```


Has And Belongs To Many (HABTM)

```
CREATE TABLE people (  
  id SERIAL PRIMARY KEY,  
  ...  
);
```

```
CREATE TABLE pets (  
  id SERIAL PRIMARY KEY,  
  ...  
)
```

```
CREATE TABLE families (  
  person_id INTEGER REFERENCES people,  
  pet_id     INTEGER REFERENCES pets  
)
```

Adding Another Owner

```
INSERT INTO people (name, age) VALUES  
( 'Danielle Mosesman', 26)
```

```
---
```

```
INSERT INTO families (person_id, pet_id) VALUES  
(1, 1),      -- John and Ruby  
(5, 1)       -- Danielle and Ruby
```

Families Table

```
SELECT *  
FROM families  
WHERE pet_id = 1      -- Ruby
```

+-----+-----+	
person_id	pet_id
+-----+-----+	
1	1
5	1
+-----+-----+	


Owners of Ruby

```
SELECT *
FROM people
  JOIN families f ON f.person_id = people.id
  JOIN pets ON pets.id = f.pet_id
WHERE pets.id = 1
```

id	name	age	person_id	pet_id	id	name	age
1	John Mosesman	26	1	1	1	Ruby	1
5	Danielle Mosesman	26	5	1	1	Ruby	1

4. Data Integrity

"Exploits of a Mom" - <https://xkcd.com/327/>



Bad Data

- Corrupted data is useless
- It makes application development more difficult

Constraints

Constraints

```
ALTER TABLE people ADD CONSTRAINT age CHECK (age > 0 AND age < 120)
```

```
---
```

```
INSERT INTO people (name, age) VALUES  
( 'Somebody' , -10)
```

```
---
```

```
ERROR:  new row for relation "people" violates check constraint "age"  
DETAIL:  Failing row contains (8, Somebody, -10).
```

Constraint Creation

```
CREATE TABLE people (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR,  
    age INTEGER CHECK (age > 0 AND age < 120)  
);
```

```
CREATE TABLE people (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR,  
    age INTEGER CONSTRAINT valid_age CHECK (age > 0 AND age < 120)  
);
```

Unique

```
CREATE TABLE people (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR,  
  age INTEGER,  
  email VARCHAR UNIQUE  
);
```

NOT NULL

```
CREATE TABLE people (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR NOT NULL,  
    age INTEGER NOT NULL  
);
```

"Tip: In most database designs the majority of columns should be marked not null."⁵

⁵ Postgres docs - <https://www.postgresql.org/docs/9.4/static/ddl-constraints.html>

Primary Key Constraint

```
CREATE TABLE people (  
    id SERIAL PRIMARY KEY  
);
```

```
CREATE TABLE people (  
    id SERIAL UNIQUE NOT NULL  
);
```

Foreign Key Constraint

```
CREATE TABLE pets (  
    id SERIAL PRIMARY KEY,  
    person_id INTEGER REFERENCES people  
)
```

```
INSERT INTO pets(person_id, name) VALUES  
(9999, 'Simba')
```

```
ERROR:  insert or update on table "pets" violates foreign key constraint "pets_person_id_fkey"  
DETAIL:  Key (person_id)=(9999) is not present in table "people".
```

=> "pets_person_id_fkey" is the default name postgres gave it

5. Performance

How do we increase performance?

- Process fewer things
- Process things faster

"Sequential Scan"

- Check every row in the table

Indexes

What is an index?

- Pre-computed data structure for efficient targeting of specific rows
- Keeps track of values and where they exist on disk
- Different types, but usually a b-tree (balanced tree)

Many Constraints Use Indexes

- UNIQUE
- PK / FK

Example

```
CREATE TABLE users (  
  name VARCHAR,  
  email VARCHAR,  
  age INTEGER,  
  created_at timestamp,  
  updated_at timestamp  
)
```

~475k rows

Made a copy - `users_2`

```
CREATE UNIQUE INDEX ON users (email)
```

Index vs No Index

```
SELECT *  
FROM users  
WHERE email = 'buck@rippin.co'
```

Un-indexed table (users_2)	- ~100ms
Indexed table (users)	- ~1ms

Q: How do we know what's actually going on?

A: View the Query Plan

Query Planner

- Most modern Database Management Systems (DBMS) have a query planner
- Query planner analyzes and optimizes queries based on internal stats it collects

EXPLAIN (No Index)

```
EXPLAIN VERBOSE SELECT *
FROM users_2
WHERE email = 'buck@rippin.co'
```

+-----+	
QUERY PLAN	
+-----+	
Seq Scan on public.users_2	(cost=0.00..11535.80 rows=1 width=63)
Output:	id, name, email, age, created_at, updated_at
Filter:	((users_2.email)::text = 'buck@rippin.co'::text)
+-----+	

EXPLAIN (With Index)

```
EXPLAIN VERBOSE SELECT *
FROM users
WHERE email = 'buck@rippin.co'
```

+-----+	
	QUERY PLAN
+-----+	
	Index Scan using users_email_idx on public.users (cost=0.42..8.44 rows=1 width=67)
	Output: id, name, email, age, created_at, updated_at
	Index Cond: ((users.email)::text = 'buck@rippin.co'::text)
+-----+	

ANALYZE

```
EXPLAIN ANALYZE VERBOSE SELECT *  
FROM users  
WHERE email = 'buck@rippin.co'
```


INDEX ALL THE THINGS!



Drawbacks of Indices

- Have to be maintained as table changes
- Take up disk space
- Lose their efficiency if they get too big
- Can hurt performance on write-heavy tables

N + 1 Queries

N + 1 Queries

```
users = User.all

users.each do |user|
  user.pets.each do |pet|
    puts pet.name
  end
end
```


N + 1 Queries

User Load (0.7ms) `SELECT "users".* FROM "users"`

Pet Load (0.4ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 1]]</code>
Pet Load (0.5ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 2]]</code>
Pet Load (0.3ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 3]]</code>
Pet Load (0.5ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 4]]</code>
Pet Load (0.4ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 5]]</code>
Pet Load (0.3ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 6]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 7]]</code>
Pet Load (0.3ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 8]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 9]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 10]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 11]]</code>
Pet Load (0.3ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 12]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 13]]</code>
Pet Load (0.2ms)	<code>SELECT "pets".* FROM "pets" WHERE "pets"."user_id" = \$1</code>	<code>[["user_id", 14]]</code>

Completed 200 OK in 290ms (Views: 266.7ms | ActiveRecord: 15.5ms)

Fixing N + 1 Queries

```
User.all => User.all.includes(:pets)
```

```
---
```

```
User Load (0.6ms)  SELECT "users".* FROM "users"
```

```
Pet Load (1.2ms)  SELECT "pets".* FROM "pets" WHERE "pets"."user_id"  
                  IN (1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14)
```

```
Completed 200 OK in 43ms (Views: 32.6ms | ActiveRecord: 8.4ms)
```

Some Final Tips For Running queries...

"Drag from the bottom up"

12
13
14
15
16
17
18
19
20
21
22
23
24
25
26

```
DELETE FROM p
  JOIN stuff ON ...
  JOIN more_stuff ON...
  AND some_condition > 50
  OR something_else = 'this'
 LEFT JOIN a_thing ON blah...
  AND thing ILIKE '%bleh%'
GROUP BY some_stuff
HAVING AVG(more_stuff) = 45
ORDER BY MAX(bleh) / MIN(bluh)
LIMIT 10
```

```
13
14
15 --DELETE FROM p
16 JOIN stuff ON ...
17 JOIN more_stuff ON...
18 AND some_condition > 50
19 OR something_else = 'this'
20 LEFT JOIN a_thing ON blah...
21 AND thing ILIKE '%bleh%'
22 GROUP BY some_stuff
23 HAVING AVG(more_stuff) = 45
24 ORDER BY MAX(bleh) / MIN(bluh)
25 LIMIT 10
26
```

14
15
16
17
18
19
20
21
22
23

```
--DELETE FROM p  
JOIN stuff ON ...  
JOIN more_stuff ON...  
AND some_condition > 50  
OR something_else = 'this'  
LEFT JOIN a_thing ON blah...  
AND thing ILIKE '%bleh%'  
GROUP BY some_stuff  
HAVING AVG(more_stuff) = 45
```



Load Query...

Save Query...

```
13
14
15 --DELETE FROM p
16 JOIN stuff ON ...
17 JOIN more_stuff ON...
18     AND some_condition > 50
19     OR something_else = 'this'
20 LEFT JOIN a_thing ON blah...
21     AND thing ILIKE '%bleh%'
22 GROUP BY some_stuff
23 HAVING AVG(more_stuff) = 45
24 ORDER BY MAX(bleh) / MIN(bluh)
25 LIMIT 10
26
```



```
12
13
14
15 --DELETE FROM p
16 JOIN stuff ON ...
17 JOIN more_stuff ON...
18     AND some_condition > 50
19     OR something_else = 'this'
20 LEFT JOIN a_thing ON blah...
21     AND thing ILIKE '%bleh%'
22 GROUP BY some_stuff
23 HAVING AVG(more_stuff) = 45
24 ORDER BY MAX(bleh) / MIN(blueh)
25 LIMIT 10
26
```



Load Query...

Save Query...

ERROR: syntax error at or near "AND"
LINE 1: AND some_condition > 50
^

11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26

```
--DELETE FROM p  
JOIN stuff ON ...  
JOIN more_stuff ON...  
    AND some_condition > 50  
    OR something_else = 'this'  
LEFT JOIN a_thing ON blah...  
    AND thing ILIKE '%bleh%'  
GROUP BY some_stuff  
HAVING AVG(more_stuff) = 45  
ORDER BY MAX(bleh) / MIN(blueh)  
LIMIT 10
```

Questions?

Thanks!

@johnmosesman